

Profiling is done using gprof and mpicc. It seems that these tools are not enough to profile MPI code and I couldn't find a proper way to do it. Here are the results:

serial

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
100.48	4.57	4.57				main
0.22	4.58	0.01	1	10.07	10.07	initMat
0.00	4.58	0.00	1	0.00	0.00	printMat

Call graph

granularity: each sample hit covers 2 byte(s) for 0.22% of 4.58 seconds

index	% time	self	children	called	name
					<spontaneous>
[1]	100.0	4.57	0.01		main [1]
		0.01	0.00	1/1	initMat [2]
		0.00	0.00	1/1	printMat [3]

		0.01	0.00	1/1	main [1]
[2]	0.2	0.01	0.00	1	initMat [2]

		0.00	0.00	1/1	main [1]
[3]	0.0	0.00	0.00	1	printMat [3]

Index by function name

[2] initMat	[1] main	[3] printMat
-------------	----------	--------------

As expected, the main computation takes up most of the time in serial implementation.

OMP

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
97.45	7.06	7.06	1	7.06	7.06	initMat
0.42	7.09	0.03	1	0.03	0.03	printMat

Call graph

granularity: each sample hit covers 2 byte(s) for 0.14% of 7.09 seconds

index	% time	self	children	called	name
					<spontaneous>
[1]	100.0	0.00	7.09		main [1]
		7.06	0.00	1/1	initMat [2]
		0.03	0.00	1/1	printMat [3]

```

          7.06  0.00  1/1      main [1]
[2]    99.6  7.06  0.00  1      initMat [2]
-----
          0.03  0.00  1/1      main [1]
[3]    0.4  0.03  0.00  1      printMat [3]
-----

```

Index by function name

[2] initMat

[3] printMat

In both MPI and OMP, the main function takes little to no time. `initMat`, the function that initializes the matrices, takes more time than the actual multiplication.

MPI

Flat profile:

Each sample counts as 0.01 seconds.

%	cumulative	self	self	total	
time	seconds	seconds	calls	Ts/call	Ts/call
100.32	1.07	1.07			main

CUDA

==75484== NVPROF is profiling process 75484, command: ./a.out 2

Time taken is 23.954443

==75484== Profiling application: ./a.out 2

==75484== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
97.70%	22.3799s	1	22.3799s	22.3799s	22.3799s	multiply(float*, float*, float*)
1.16%	265.92ms	1	265.92ms	265.92ms	265.92ms	[CUDA memcpy DtoH]
1.14%	261.20ms	2	130.60ms	130.58ms	130.62ms	[CUDA memcpy HtoD]

==75484== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
99.55%	22.9089s	3	7.63629s	130.68ms	22.6470s	cudaMemcpy
0.44%	101.02ms	3	33.673ms	419.65us	100.17ms	cudaMalloc
0.00%	914.48us	3	304.83us	271.89us	366.54us	cudaFree
0.00%	659.10us	83	7.9400us	502ns	311.76us	cuDeviceGetAttribute
0.00%	58.119us	1	58.119us	58.119us	58.119us	cudaLaunch
0.00%	48.038us	1	48.038us	48.038us	48.038us	cuDeviceTotalMem
0.00%	41.055us	1	41.055us	41.055us	41.055us	cuDeviceGetName
0.00%	9.2000us	3	3.0660us	315ns	7.9910us	cudaSetupArgument
0.00%	3.9460us	1	3.9460us	3.9460us	3.9460us	cudaConfigureCall
0.00%	3.5950us	2	1.7970us	701ns	2.8940us	cuDeviceGetCount
0.00%	1.7160us	2	858ns	727ns	989ns	cuDeviceGet
0.00%	494ns	1	494ns	494ns	494ns	cudaGetLastError

Here memory transfer takes considerable amount of time apart from the time taken by `multiply`. However, for smaller matrix sizes `memcpy` takes more time than `multiply`.

Cachegrind analysis: The instruction miss rates and data miss rates are mentioned below. We can see that serial and OMP has similar statistics where as in MPI instruction and data-miss rates are considerably higher.

serial

```

==2927== Cachegrind, a cache and branch-prediction profiler
==2927== Copyright (C) 2002-2015, and GNU GPL'd, by Nicholas Nethercote et al.
==2927== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==2927== Command: ./serial
==2927==
--2927-- warning: L3 cache found, using its data for the LL simulation.

```

```

Time taken is 235.619629
==2927==
==2927== I   refs:      33,037,538,463
==2927== I1  misses:      1,819
==2927== LLi misses:      1,690
==2927== I1  miss rate:      0.00%
==2927== LLi miss rate:      0.00%
==2927==
==2927== D   refs:      14,019,014,062 (13,015,728,675 rd + 1,003,285,387 wr)
==2927== D1  misses:      62,772,493 ( 62,644,822 rd + 127,671 wr)
==2927== LLd misses:      62,761,087 ( 62,633,938 rd + 127,149 wr)
==2927== D1  miss rate:      0.4% ( 0.5% + 0.0% )
==2927== LLd miss rate:      0.4% ( 0.5% + 0.0% )
==2927==
==2927== LL refs:      62,774,312 ( 62,646,641 rd + 127,671 wr)
==2927== LL misses:      62,762,777 ( 62,635,628 rd + 127,149 wr)
==2927== LL miss rate:      0.1% ( 0.1% + 0.0% )

```

OMP

```

==3133== Cachegrind, a cache and branch-prediction profiler
==3133== Copyright (C) 2002-2015, and GNU GPL'd, by Nicholas Nethercote et al.
==3133== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==3133== Command: ./omp
==3133==
--3133-- warning: L3 cache found, using its data for the LL simulation.

```

```

Time taken is 244.495880
==3133==
==3133== I   refs:      33,037,538,399
==3133== I1  misses:      1,816
==3133== LLi misses:      1,687
==3133== I1  miss rate:      0.00%
==3133== LLi miss rate:      0.00%
==3133==
==3133== D   refs:      14,019,014,043 (13,015,728,664 rd + 1,003,285,379 wr)
==3133== D1  misses:      62,772,305 ( 62,644,708 rd + 127,597 wr)
==3133== LLd misses:      62,761,087 ( 62,633,937 rd + 127,150 wr)
==3133== D1  miss rate:      0.4% ( 0.5% + 0.0% )
==3133== LLd miss rate:      0.4% ( 0.5% + 0.0% )
==3133==
==3133== LL refs:      62,774,121 ( 62,646,524 rd + 127,597 wr)
==3133== LL misses:      62,762,774 ( 62,635,624 rd + 127,150 wr)
==3133== LL miss rate:      0.1% ( 0.1% + 0.0% )

```

MPI

```

==3225== Cachegrind, a cache and branch-prediction profiler
==3225== Copyright (C) 2002-2015, and GNU GPL'd, by Nicholas Nethercote et al.
==3225== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==3225== Command: mpirun -np 8 ./mpi
==3225==

```

```

--3225-- warning: L3 cache found, using its data for the LL simulation.
[warn] Epoll ADD(4) on fd 1 failed. Old events were 0;
read change was 0 (none); write change was 1 (add): Operation not permitted

```

```

Time taken is 3.978104
==3225==
==3225== I   refs:      60,582,190
==3225== I1  misses:      83,243
==3225== LLi misses:     10,336
==3225== I1  miss rate:      0.14%
==3225== LLi miss rate:      0.02%
==3225==
==3225== D   refs:      24,637,760 (16,965,706 rd + 7,672,054 wr)
==3225== D1  misses:      849,555 ( 795,670 rd +   53,885 wr)
==3225== LLd misses:      76,610 ( 48,015 rd +   28,595 wr)
==3225== D1  miss rate:      3.4% ( 4.7% +   0.7% )
==3225== LLd miss rate:      0.3% ( 0.3% +   0.4% )
==3225==
==3225== LL refs:      932,798 ( 878,913 rd +   53,885 wr)
==3225== LL misses:      86,946 ( 58,351 rd +   28,595 wr)
==3225== LL miss rate:      0.1% ( 0.1% +   0.4% )

```